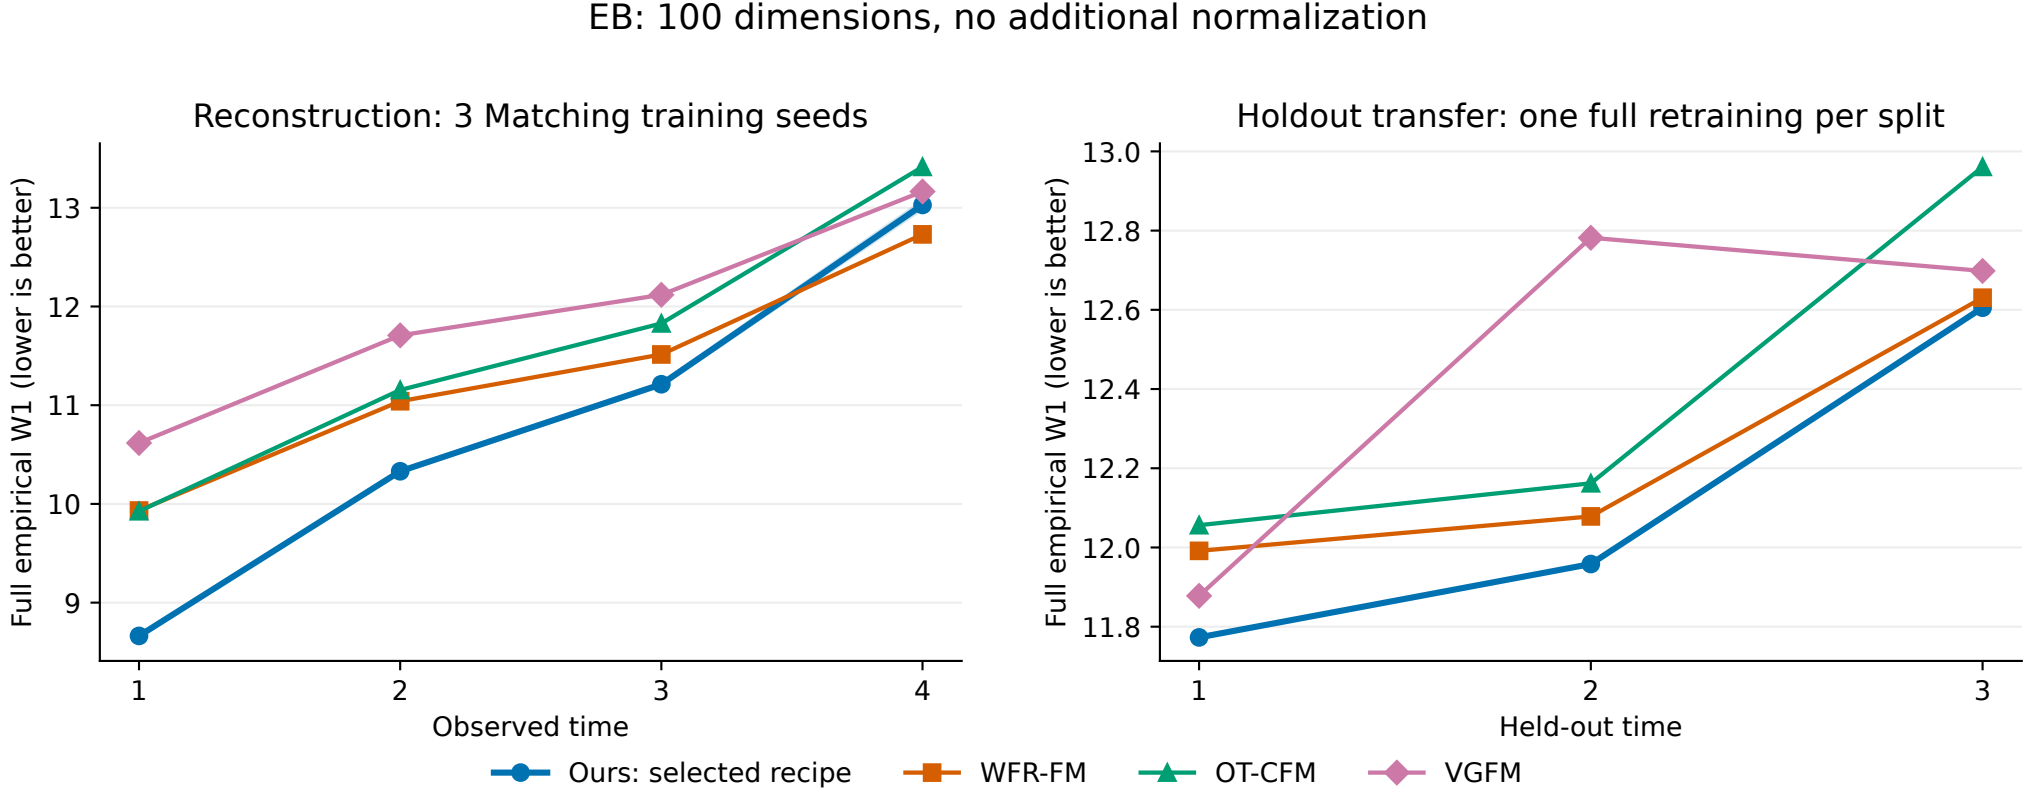
EB: 100 dimensions, no additional normalization
Reconstruction: 3 Matching training seeds
Holdout transfer: one full retraining per split
Full empirical W1 (lower is better)
Observed time
Held-out time
Ours: selected recipe
WFR-FM
OT-CFM
VGFM
Shading: ±1 SD across Matching seeds, with a fixed path and couplings. Baselines are archived runs.
The recipe was selected on reconstruction before evaluating holdout predictions; budgets are not matched.